

LIVING AI SOLUTION DOSSIER

Illustrative Sample Excerpt

The reviewed artifact at the center of the program.

FRAME · DESIGN · PROVE · FORESEE

PARTICIPANT	Maya R. <i>(fictional)</i>	FIELD	Professional services / advisory ops
ROLE	Senior Client Delivery & Operations Lead	ORGANIZATION	Mid-sized advisory firm (anonymized)
DOSSIER STATUS	Illustrative — pre-final review stage	METHOD	The TenX Method

This is an illustrative sample. It is not a real participant dossier and contains no client-confidential information. “Maya R.” is a fictional persona created to show the structure, depth, and review standard of a TenXPros Living AI Solution Dossier. Names, figures, and examples are invented for demonstration. The review labels are illustrative and are not a certification decision.

You bring the expertise. We bring the AI method. Together, your professional work becomes TenX.



ORIENTATION

What this sample shows

Five things to hold in mind as you read. This excerpt is built to demonstrate a standard, not to make a claim.

- **It's an illustrative sample, not a real participant case.** Maya R. is fictional. Nothing here represents a real person, client, or outcome.
- **It shows the thinking, structure, and evidence we expect.** This is the *kind* of work a TenXPros dossier contains — judgment made explicit, not notes or a quiz score.
- **It includes reviewer notes, because review is the product.** You're seeing how work is assessed, not just what it looks like. The credential is earned through that review.
- **It intentionally leaves one evidence gap.** The value case is built but not yet measured — included on purpose, to show exactly where the certification bar sits.
- **A fully certifiable dossier closes that gap with measured pilot evidence.** This sample sits one honest step short of certification, and the final pages show precisely what that step is.

- Throughout this dossier, one boundary holds: **AI drafts; a named human reviews and owns every client-facing output.**

EXECUTIVE SNAPSHOT

The case at a glance

PROFESSIONAL CONTEXT. Maya leads client delivery and operations at a mid-sized advisory firm. She owns the path from first client conversation to signed proposal to the handoff that begins delivery. Three to four engagement teams rely on her to keep that path consistent.

THE PROBLEM. Discovery, proposal drafting, and handoff are inconsistent, slow, and dependent on individual memory. Quality varies with who staffs the work. Context captured in discovery is often lost by delivery, forcing teams to re-interview clients.

WHY AI IS RELEVANT. The lost value sits in *unstructured language work* — notes, briefs, draft sections, summaries — where AI drafting genuinely helps. The opportunity is a responsible, human-reviewed workflow, not a system that makes client judgments.

WHAT AI WILL DO. Draft first-pass discovery summaries from approved notes; assemble proposal sections from a vetted library; produce structured handoff briefs; flag missing information.

WHAT AI WILL NOT DO. Decide whether to pursue a client, set pricing, make commitments, or replace judgment on scope and risk. **AI drafts; a named human reviews and owns every client-facing output.**

EXPECTED VALUE (HYPOTHESIS, NOT GUARANTEE). More consistent discovery and proposals; less context lost at handoff; faster drafting; a clearer audit trail. Each must be measured against a baseline before it is treated as proven.

- Illustrative review outcome:** a strong draft, certifiable after evidence revision. The design and governance are sound; the one open item is measured pilot evidence (see *What would make this fully certifiable?*).

DIMENSION	ILLUSTRATIVE LABEL
Problem definition	Strong
Risk & boundaries	Strong
Use-case selection	Strong
Evaluation rubric & test set	Solid — quantify thresholds
Workflow usability	Solid — one gap to close
Value evidence	Not yet measured (the intentional gap)
Governance & confidentiality	Strong
90-day roadmap	Solid — add numeric gates

90-DAY DIRECTION. A narrow, measured pilot on one use case (discovery summary), with a human review gate at every step, before any expansion.

The dossier at a glance

A complete Living AI Solution Dossier has **twelve sections**, produced across the four phases of the TenX Method. This excerpt presents the full structure and shows the key sections in depth; lighter sections are summarized to keep the sample readable.

#	SECTION	PHASE
1	Professional Context	Frame
2	Problem Definition	Frame
3	AI Suitability Assessment	Frame
4	Context, Stakeholder & Initial Foresight Analysis	Frame
5	Data & Evidence Review	Design
6	Workflow Before / After	Design
7	Risk, Ethics, Privacy & Compliance Review	Design
8	Responsible AI Solution Design	Design
9	Adoption & Communication Plan	Prove
10	Value, Roadmap & Proof Plan	Prove
11	Personal AI Foresight Plan	Foresee
12	Final Recommendation	Foresee

MAPPING NOTE

This excerpt follows the canonical twelve-section structure above. In an earlier instructional draft, the use-case portfolio, grounded-knowledge pack, and evaluation rubric appeared as standalone sections; here they sit inside their canonical homes — the use-case portfolio under §4, the knowledge pack under §5, and the rubric and test set under §10 (the proof plan). The content is the same; the labels now match the dossier’s official anatomy.

SECTION 1

Professional Context

PHASE · FRAME

Maya is accountable for the operational quality of client engagements from first contact to the start of delivery. In practice, three responsibilities sit with her:

- **Discovery integrity** — making sure the firm understands a client before proposing work, and that what’s learned is written down usefully.
- **Proposal consistency** — ensuring proposals reflect the firm’s standards and prior experience, not whoever happens to be free.
- **Handoff clarity** — transferring context so delivery starts informed, not from zero.

The work is **language-heavy and judgment-heavy**. Most of what she produces is written, but the value is in knowing what matters, what to ask, and what to leave out. Her constraint isn’t effort; it’s consistency and memory. When she’s involved, quality is high. When she’s not, it varies.

♦ REVIEWER NOTE – STRONG

Maya separates **language work** (AI-suitable) from **judgment work** (not) before naming a single tool. Weak dossiers describe a job title; strong ones describe the decisions and failure modes inside the work.

SECTION 2

Problem Definition

PHASE · FRAME

PROBLEM STATEMENT.

Across discovery, proposal drafting, and handoff, quality depends on which individual is involved. Knowledge captured early is lost later. The result is rework, inconsistent proposals, and engagements that begin under-informed — a cost the firm absorbs invisibly.

CURRENT PAIN.

- Discovery notes live in personal formats and are rarely reusable.
- Proposals are rebuilt each time; strong past language isn’t systematically reused.
- Handoff is often a short verbal summary, so delivery re-interviews clients.
- The firm can’t easily see why a past proposal was structured as it was.

OUT OF SCOPE (DELIBERATELY). Client selection, pricing, and scope decisions; any confidential data entering an unvetted tool; replacing human review of client-facing work; automating the relationship or the judgment.

WHY IT’S WORTH SOLVING. The cost is recurring and compounding, and it lives in exactly the structured language work a reviewed AI workflow can help — without touching client judgment. Solving it reduces key-person dependency, an operational risk.

WHAT “GOOD” LOOKS LIKE. Discovery captured in a consistent, reusable structure; proposals that start from vetted components; handoff briefs complete enough to start delivery informed; a human reviewing and owning every client-facing output; a visible trail.

REVIEWER NOTE – STRONG, WITH ONE FIX

Clear, bounded problem, and the out-of-scope list does real protective work. **To certify:** at least one dimension of “good” (e.g. context lost at handoff) needs a measurable definition so success can be proven, not asserted.

SECTION 3

AI DRAFTS · A HUMAN OWNS THE OUTPUT

AI Suitability Assessment

PHASE · FRAME

WHERE AI HELPS. The bottleneck is repeated, unstructured language work — summarizing notes, assembling sections, producing briefs. Language models are well-suited to *first drafts* and *structured extraction* from approved material: the low-judgment, high-volume layer of Maya’s work.

WHERE AI SHOULD NOT BE TRUSTED. Client judgment (whether to pursue, what to charge, what to promise); factual claims about a client not traceable to an approved source; final client-facing language without review; anything implying a commitment.

WHERE HUMAN JUDGMENT STAYS. Every client-facing artifact passes through a named, accountable human. **AI produces a draft; a person owns the output.**

CONFIDENTIALITY & DATA BOUNDARIES. No confidential names or identifiers enter the workflow without redaction; only approved, firm-owned sources are used; sensitive engagements are excluded; the posture is moderate sensitivity, designed for redacted inputs and human verification.

SUITABILITY JUDGMENT. **Suitable, with guardrails.** A strong fit for AI *assistance* at the drafting layer; a poor fit for AI *decision-making*. Suitability holds only while human review and data boundaries are enforced.

♦ REVIEWER NOTE – STRONG

Conditional suitability (“suitable *while* these guardrails hold”) is more mature than a numeric score, because it names what would invalidate the assessment. **Open question:** name the specific control that stops a confidential name entering the tool — a redaction step, a checklist, or a filter.

SECTION 4

AI DRAFTS · A HUMAN OWNS THE OUTPUT

Context, Stakeholder & Initial Foresight Analysis

PHASE · FRAME

WHO IS AFFECTED.

STAKEHOLDER	HOW THE PROBLEM AFFECTS THEM
Clients	Repeated questions; slower starts; uneven experience
Delivery teams	Begin under-informed; absorb rework
Partners	Inconsistent proposal quality reaching review

STAKEHOLDER	HOW THE PROBLEM AFFECTS THEM
Maya	A personal bottleneck; quality drops when unavailable
The firm	Invisible cost in time, rework, inconsistency

THE AI OPPORTUNITY LANDSCAPE (USE-CASE PORTFOLIO) . Maya scored seven candidates on **value**, **risk**, **feasibility**, and **time-to-value** (1–5; 5 = most favorable, i.e. high value, low risk, high feasibility, fast).

#	USE CASE	VALUE	RISK (5=LOW)	FEAS.	TIME-TO-VALUE	NOTE
1	Discovery summary from approved notes	4	4	5	5	Primary. Low risk, high frequency, clear review point
2	Proposal section assembly (vetted library)	5	3	3	3	High value; needs a curated library first
3	Delivery handoff brief generation	4	4	4	4	Strong phase two once #1 is reliable
4	Missing-information flagging	3	4	4	4	Useful guardrail; pairs with #1
5	Meeting-notes structuring	3	4	5	5	Easy, lower standalone value
6	Proposal consistency checklist	3	3	3	2	Valuable later; needs a defined standard
7	Past-engagement search	4	2	2	2	High value, higher risk; access control. Deferred

Scores are illustrative judgments, not benchmarks. 5 = most favorable on each axis.

PRIMARY USE CASE. #1 — **Discovery summary from approved notes.**

WHY THIS FIRST. Highest feasibility, lowest risk, fastest to value, frequent enough to measure quickly, and it produces the structured input later use cases (#3, #4) depend on.

INITIAL FORESIGHT. If #1 proves reliable, the natural sequence is #4 (flagging) → #3 (handoff briefs) → #2 (proposal assembly). Each step expands only on evidence from the last.

♦ REVIEWER NOTE – STRONG

Correct sequencing. Maya chose the highest-*confidence* use case over the highest-*value* one (#2), and deferred the riskiest (#7) on access-control grounds. This prioritization judgment is exactly what the program rewards.

Data & Evidence Review

PHASE · DESIGN

The workflow must be **grounded** — it works only from approved material, never from open-ended guesses about a client.

APPROVED SOURCES. Maya’s discovery notes (redacted of identifiers); the firm’s standard question set; the engagement taxonomy; anonymized structural patterns from past summaries (format, not content).

REUSABLE, FIRM-OWNED CONTEXT. The standard discovery template; service-line definitions; an internal glossary so summaries use consistent language.

FORBIDDEN DATA (NEVER ENTERS). Confidential names, logos, or identifiers (pre-redaction); pricing, contracts, commercial terms; third-party confidential material; unnecessary personal data.

SOURCE QUALITY RULES.

- Every factual statement must trace to an approved source.
- If a required field has no source, the output marks it “**open question**” — it never invents an answer.
- Redaction happens *before* material enters the workflow; a human confirms it.

ILLUSTRATIVE SNIPPET — INPUT → GROUNDED OUTPUT (REDACTED)

INPUT “[CLIENT], a [SECTOR] org with ~[N] staff. Concern: [ISSUE], handled manually. Wants [OUTCOME] within [TIMEFRAME]. Constraints: limited capacity; a prior rollout stalled.”

OUTPUT “**Situation:** a mid-sized [SECTOR] organization seeks to improve [PROCESS]. **Goal:** [OUTCOME] within [TIMEFRAME]. **Constraints:** limited capacity; one prior stalled rollout. **Open questions:** budget authority; success metric; appetite for change. (Draft — requires human review.)”

◆ REVIEWER NOTE — STRONG

The rule “if no source, mark as open question, never invent” converts the model’s biggest weakness — confident fabrication — into a visible, reviewable signal. Redaction-before-entry is the correct sequence.

SECTION 6

AI DRAFTS · A HUMAN OWNS THE OUTPUT

Workflow Before / After

PHASE · DESIGN

CURRENT WORKFLOW.

STEP	WHO	TODAY	FAILURE POINT
Discovery call	Maya / lead	Personal-format notes	Inconsistent capture
Summary	Maya	Written if time allows	Often skipped
Proposal draft	Whoever staffed	Rebuilt from memory	Quality varies
Handoff	Maya → delivery	Short verbal summary	Context lost

AI-ASSISTED WORKFLOW (PRIMARY USE CASE).

STEP	WHO	REVIEW CHECKPOINT
Discovery call	Maya / lead	—
Redaction	Lead	Human confirms redaction
Draft summary	AI	—
Review & approve	Maya	Human owns the output
Feeds handoff	Maya → delivery	Human confirms completeness

WHAT IS NOT AUTOMATED. The discovery conversation, the judgment about what matters, the decision to pursue work, and final approval of any client-facing artifact. The workflow speeds the *writing around* the judgment, not the judgment.

REVIEWER NOTE — SOLID, ONE GAP

Checkpoints sit at the right moments. **To certify:** specify *who* performs the final completeness check when Maya is unavailable — otherwise the key-person dependency quietly returns at handoff.

SECTION 7

AI DRAFTS · A HUMAN OWNS THE OUTPUT

Risk, Ethics, Privacy & Compliance Review

PHASE · DESIGN

RISK	CONCERN	MITIGATION
Data privacy	Client data in an un-vetted tool	Redaction before entry; approved sources only; sensitive work excluded
Hallucination	Invented client facts	Grounding rule (“open question,” never invented); faithfulness check
Over-reliance	Drafts treated as final	Outputs labeled “draft — requires review”; human owns every output
Confidentiality	Identifiers leak	Mandatory redaction-confirm step; test case gates the pilot
Bias / inconsistency	Summaries skew or vary	Standard template + rubric; periodic spot-checks
Accountability	Unclear ownership	Named approver per artifact; approval recorded

ACCOUNTABILITY PRINCIPLE. AI is a drafting assistant. A named person is accountable for every client-facing output, exactly as today. **Accountability does not move to a tool.**

◆ REVIEWER NOTE — STRONG

Governance is designed in, not bolted on. This is where many dossiers overclaim safety; this one names specific, checkable controls instead.

Responsible AI Solution Design

PHASE · DESIGN

The designed solution is a **grounded, human-reviewed assistant** built around one recurring task.

DESIGN PRINCIPLES.

- **Grounded:** the assistant answers only from approved notes and templates; unsourced fields become open questions.
- **Labeled draft:** every output is explicitly a draft, never a final.
- **Reusable interaction pattern:** a fixed structure — capture → redact → draft → review → approve — applied the same way each time, so the team builds a shared habit, not ad-hoc prompting.
- **Human-owned:** a named reviewer approves before anything moves downstream.

FAILURE POINTS THE DESIGN ANTICIPATES. Fabrication (mitigated by grounding); redaction slips (mitigated by a confirm step); over-trust (mitigated by draft labeling); silent drift (mitigated by rubric spot-checks).

REVIEWER NOTE – SOLID

The reusable capture → redact → draft → review → approve pattern is the difference between a tool and a *workflow*. **To strengthen:** document the assistant’s instructions as a reusable artifact so a second person could run it identically.

Adoption & Communication Plan

PHASE · PROVE

A workflow only creates value if the teams actually use it — correctly.

- **Who leads it.** Maya owns the pilot; one delivery lead acts as a second reviewer so the practice isn’t dependent on her alone.
- **How it’s introduced.** A short working session for the three to four teams: what the assistant does, what it does *not* do, and the one rule that matters — *it’s a draft; you own the output*.
- **What’s communicated.** The redaction step is mandatory; “open questions” are normal and good; nothing client-facing skips review.
- **How feedback flows.** Reviewers log any weak output against the rubric, so issues are visible and fixable rather than quietly worked around.

REVIEWER NOTE – ADEQUATE

Sensible and lightweight. **To strengthen:** name how adoption itself is observed — e.g. the share of summaries produced through the workflow versus written from scratch.

Value, Roadmap & Proof Plan

PHASE · PROVE

This is the heart of the evidence — and the section that shows where certification standards bite.

PROOF PLAN – EVALUATION RUBRIC.

CRITERION	“PASS” MEANS
Faithfulness	Every statement traces to the notes; nothing invented
Completeness	All template fields present, or marked “open question”
No fabricated specifics	No invented facts, names, numbers, or commitments
Appropriate uncertainty	Gaps surfaced honestly
Usable structure	A delivery lead could act on it without rewriting
Confidentiality	No un-redacted identifiers

PROOF PLAN – TEST SET (5 CASES).

#	INPUT	EXPECTED BEHAVIOR	RISK IF FAILED
1	Clean, complete notes	Accurate summary, no gaps	Low
2	Missing budget field	Marks budget “open question”; no invented figure	High
3	Client name (redaction missed)	Flags or omits the identifier	High
4	Ambiguous / contradictory notes	Surfaces the contradiction	Med
5	Sparse notes	Short summary; flags what’s missing	Med

Cases 2 and 3 must pass every time before the pilot is trusted; a failure halts expansion.

VALUE CASE – HYPOTHESES, NOT RESULTS.

No value below is claimed as achieved. Each is a hypothesis with a measurement.

HYPOTHESIS	HOW IT’S MEASURED	HONEST FRAMING
Time saved on summaries	Net time per summary, before vs. during pilot	Real only after review time is included
More consistent proposals	Rubric-score variance across authors	Consistency is the goal; speed is secondary
Less rework at handoff	Count of re-interviews / clarifications	Hardest to attribute; needs a clean baseline
Faster handoff	Days from discovery to informed start	Confounded by client factors; read cautiously
Lower key-person dependency	Quality when Maya is <i>not</i> involved	The real prize; slow to observe

90-DAY ROADMAP.

PHASE	FOCUS	SUCCESS CRITERIA	REVIEW GATE
Days 1–30	Baseline & build	Baselines recorded; passes test cases 2 & 3 in trials	Gate 1: no expansion until high-risk cases pass
Days 31–60	Narrow pilot	≥ agreed pass rate on rubric; zero confidentiality incidents	Gate 2: halt on any case-2/3 failure or incident
Days 61–90	Measure & decide	A measured before/after on ≥1 metric; documented decision	Gate 3: expand only if value <i>and</i> safety hold

♦ REVIEWER NOTE – THE DECISIVE SECTION

The proof plan is well-targeted: the two high-risk test cases protect against the failures that would actually cause harm, and the gates have teeth. **The intentional gap lives here:** there is no baseline yet, so value cannot be *shown*. As written, this is a strong plan to prove value — not yet proof of value. Closing that is the step from “strong draft” to “certified” (see next page).

Personal AI Foresight Plan

PHASE · FORESEE

Beyond the first workflow, Maya’s role is to lead how AI enters her function responsibly.

- **The strategic aim.** Reduce key-person dependency: make discovery quality a property of the *system*, not of who is in the room.
- **What she’ll govern.** The redaction discipline, the review standard, and the decision to expand — each tied to evidence, not enthusiasm.
- **What she’s learning to watch.** Over-trust in drafts; quiet drift in quality; any erosion of the “human owns it” boundary.
- **How she’ll decide what’s next.** Expansion to handoff briefs and proposal assembly is considered only after the pilot clears its gates.

REVIEWER NOTE — STRONG DIRECTION

This reframes the participant from *tool user* to *adoption leader* — the program’s actual outcome. **To strengthen:** add one forward indicator she’ll review quarterly (e.g. summary quality on engagements she didn’t touch).

SECTION 12

AI DRAFTS · A HUMAN OWNS THE OUTPUT

Final Recommendation

PHASE · FORESEE

- **Recommendation: PILOT** — proceed to a narrow, measured pilot, not a full rollout.

WHY. The primary use case is high-value, low-risk, frequent, and reviewable. The design turns the model’s main weakness into a visible signal, governance is concrete, and value is measured honestly rather than assumed.

CONDITIONS. Capture a baseline *before* any change; high-risk test cases pass before the pilot is trusted; every client-facing output is human-reviewed and owned; any confidentiality incident halts the pilot.

WHAT STAYS HUMAN-LED, PERMANENTLY. The discovery conversation; the judgment about what matters; client-pursuit, scope, and pricing; and final approval of every client-facing artifact. **The workflow assists the writing around these decisions — it never makes them.**

◆ REVIEWER SUMMARY (ILLUSTRATIVE)

A disciplined, defensible plan with the right first use case and honest economics. **Illustrative review outcome: a strong draft, certifiable after evidence revision** — the work is sound; what remains is measured pilot evidence and a few quantified thresholds. This is a *pre-final* review state, shown deliberately so you can see where the bar sits. (*Illustrative only; not an actual decision.*)

THE REVIEW STANDARD

AI DRAFTS · A HUMAN OWNS THE OUTPUT

What would make this fully certifiable?

This sample sits one honest step short of certification. The step is **evidence**. To move from *strong draft* to *certified*, the dossier would add:

ITEM	WHAT’S NEEDED
Baseline measurement	Capture each metric (net time, rework, handoff days) <i>before</i> the pilot changes anything.
Pilot sample size	A stated number of summaries large enough that a difference isn’t noise (e.g. a defined N across the 30–60 day window).
Pass-rate threshold	A named target for the rubric (e.g. high-risk cases 2 & 3 at 100%; cases 1/4/5 at an agreed rate).
Before/after comparison	At least one metric measured both before and after, reporting <i>net</i> effect (drafting time saved minus review time added).
Named handoff owner	A specific second reviewer who performs the completeness check when Maya is unavailable.
Quantified review gates	Each 90-day gate stated as a number, so it’s an objective stop/go — not a judgment call in the moment.

Close these six, and the same dossier reads as **certified**: strong thinking *plus* evidence it held up.

REFERENCE

AI DRAFTS · A HUMAN OWNS THE OUTPUT

Eight Assets Map

A Living AI Solution Dossier assembles the eight assets every participant builds. Here is where each appears in this sample, and why it matters professionally.

ASSET	WHERE IT APPEARS	WHY IT MATTERS
Personal AI Strategy Brief	§1–3 — where AI adds value, where it’s dangerous	The judgment layer: deciding where AI belongs before touching a tool

ASSET	WHERE IT APPEARS	WHY IT MATTERS
AI Use-Case Portfolio	§4 — seven candidates, ranked	Prioritizing under value, risk, and feasibility is the core adoption skill
Interaction & Decision Kit	§5 question set + §8 capture → review pattern	Reusable patterns that make outputs consistent and decisions sharper
Grounded Domain Knowledge Pack	§5 — approved sources, forbidden data, rules	Makes your expertise safe and reusable for AI
AI Evaluation Rubric & Test Set	§10 — rubric + five test cases	How you measure quality and catch failure modes; the basis of trust
Custom Assistants & AI Workflows	§6 + §8 — the designed, grounded workflow	A working tool built around a real, recurring task
AI Value & Economics Case	§10 — hypotheses + measurement plan	Impact stated in terms a stakeholder will accept
Final Portfolio & 90-Day Roadmap	§10 roadmap + §12 + the dossier as a whole	The defended artifact and a credible plan for what’s next

Rubric Mapping — the TenXPros 8 review criteria



The credential is earned through reviewed evidence — not attendance.

How this illustrative dossier maps to the public review standard. Labels are illustrative.

#	CRITERION	LABEL	WHY
1	The problem is clearly defined	Strong	Specific, bounded, with stakeholders and explicit out-of-scope lines (§2, §4)
2	Risks and boundaries are explicit	Strong	Named risks with mitigations; “what AI will not do” stated and held (§3, §7, §12)
3	Use-cases are correctly chosen and prioritized	Strong	Scored portfolio; highest-confidence use case first; risky retrieval deferred (§4)
4	An evaluation rubric and test set exist	Adequate	Rubric and five well-targeted test cases present; thresholds and sample size to be quantified (§10)
5	The workflow is genuinely usable	Adequate	Clear before/after with review gates; one coverage gap — handoff owner when Maya is away (§6)
6	Value and impact are shown with evidence	Needs revision	Honest hypotheses and a measurement plan, but no baseline captured yet — the intentional gap (§10)
7	Governance and confidentiality are respected	Strong	Redaction-before-entry, grounded sources, named accountability, incident-halt rule (§5, §7)
8	The 90-day roadmap is realistic	Adequate	Phased, gated, cautious; gates need numeric thresholds to be objective (§10)

Overall (illustrative): strong design and governance; the path to full certification runs through **evidence**. This is precisely the difference between *attendance* and *reviewed work* — the dossier is judged on whether the thinking holds up, not on whether the program was finished.

DISCLAIMER

This is an illustrative sample created to demonstrate the structure and review standard of a TenXPros Living AI Solution Dossier. “Maya R.” is fictional. No real client, person, organization, or confidential information is represented. Review labels are illustrative and do not constitute a certification decision.

The TenX Method: Frame · Design · Prove · Foresee — You bring the expertise. We bring the AI method.